

STG17 TECHNICAL WORKSHOP · DAY 3 · WEDNESDAY AFTERNOON

Engines of Scale

Searching, aggregating and mapping big data
with the Elastic Stack (ELK)

*From Ookla Speedtest tiles to a population-weighted,
subnational connectivity indicator*

Emerging Issues, Emerging Practice — Innovating the Data Value Chain

African Development Bank · African Union (STATAFRIC) · National Institute of Statistics of Rwanda

This afternoon at a glance

14:00 – 14:30

TALK



Engines of Scale

Big data building blocks, how the Elastic Stack works, and how to use it responsibly in an NSO.

14:45 – 15:30

LAB 1



Oookla at scale with Elasticsearch

Create the mapping, index the tiles, run the first aggregations and geospatial queries.

15:45 – 16:45

LAB 2



Search-driven exploration

Kibana dashboard by region and quarter, benchmark against pandas, publish on GitHub.

SIX PARTS OF THE TALK

01

Why scale matters

From laptop to cluster: when size changes the tools

02

The Elastic Stack explained

E, L and K — and the inverted index behind them

03

Modelling statistical data

Documents, mappings and geo types for Oookla tiles

04

Querying and aggregating

Query DSL, aggregations, geo queries and ES|QL

05

Exploring with Kibana

Discover, Lens, Maps and the pandas benchmark

06

Governing and sustaining

Quality, licensing, deployment and publication

What you will be able to do by 17:00



Explain

where indexing and search sit in the big data value chain, next to lakes, lakehouses and Spark.



Describe

the Elastic Stack components and the anatomy of a cluster: nodes, indices, shards, replicas.



Design

an explicit mapping for Ookla tiles, including keyword, numeric and geo_shape fields.



Query

with filters, aggregations and geospatial queries — in Query DSL and in ES|QL.



Explore

connectivity by region and quarter in a Kibana dashboard that non-coders can use.



Judge

with your own measurements whether Elasticsearch is worth its complexity for your office.

Feeds the STG17 Action Plan 2025–2030: 2.1.1 Guidelines on new data sources · 4.2 Capacity building · Day 5 country results

One session, four profiles — what each of you should take away



Statisticians & methodologists

YOUR FOCUS

Weighting, coverage bias and approximate aggregations — what makes a figure publishable, or not.



Data scientists & analysts

YOUR FOCUS

Mappings, bulk indexing, Query DSL and ES|QL — reusable code for the team repository.



IT & infrastructure staff

YOUR FOCUS

Cluster anatomy, shard sizing, security and deployment options under sovereignty constraints.



Managers & decision-makers

YOUR FOCUS

When the investment is justified, what it costs in skills, and how it serves the Action Plan.



No prerequisites in Elasticsearch. The cluster is pre-provisioned, notebooks are ready, and Kibana lets non-coders explore. Mix profiles in each team: one coder, one methodologist.

SECTION 01


01

Why scale matters


When does a laptop and pandas stop being enough?



The big data value chain: six building blocks



CROSS-CUTTING



Lake, lakehouse or search index: which does what?



Data lake

Answers: Where is the raw file?

- Files on object storage (Parquet, CSV, rasters)
- Schema-on-read, very low cost
- Keeps the untouched source for audit



Lakehouse

Answers: What is the official aggregate?

- Tables on the lake (Delta Lake, Apache Iceberg)
- SQL, transactions and time travel
- Reproducible production of statistics



Search index

Answers: What does this slice look like, now?

- Elasticsearch / OpenSearch
- Inverted index and BKD trees
- Sub-second filters, text, geo, dashboards



Batch

Process a bounded dataset at intervals — e.g. each quarterly Ookla release.



Streaming

Process unbounded events as they arrive — e.g. sensors, network or transaction feeds.

The Ookla case: why this dataset tests your tools

610.8 m

tile edge at the equator (web-Mercator
zoom level 16)

Q1 2019

first quarter available; new tiles every
quarter since

2 series

fixed broadband and mobile, published
separately

KEY FIELDS IN EACH TILE

quadkey

avg_d_kbps

avg_u_kbps

avg_lat_ms

tests

devices

tile (WKT)



Why a single laptop struggles. Ookla aggregates hundreds of millions of monthly tests into tiles, yet a multi-country, multi-quarter, fixed + mobile panel still grows fast. Add geometry, joins to boundaries and WorldPop, and several colleagues asking questions at the same time — and memory, not skill, becomes the bottleneck.

Source: Ookla, [ookla-open-data README \(GitHub\)](#) and [Registry of Open Data on AWS](#) — licence CC BY-NC-SA 4.0.

SECTION 02


02

The Elastic Stack explained

What exactly are E, L and K — and how do they work together?



ELK in one picture



From "ELK" to "Elastic Stack". Beats (lightweight data shippers) joined E, L and K, so Elastic renamed the family. People still say "ELK" — both names refer to the same toolset.

The core idea: the inverted index

DOCUMENTS

D1	Kigali · urban · mobile
D2	Musanze · rural · fixed
D3	Kigali · rural · mobile

index once

TERM → DOCUMENTS

kigali	D1, D3
musanze	D2
urban	D1
rural	D2, D3
mobile	D1, D3
fixed	D2

QUERY

kigali AND rural

$\{D1, D3\} \cap \{D2, D3\}$
= {D3}

Found without reading a single document — only two short lists are compared.




Think of the index at the back of a statistical yearbook: look up the term, jump straight to the pages. Numbers and coordinates use a sibling structure (BKD trees) for ranges and geography. The trade-off: indexing costs time, disk and memory up front, so queries can be fast later.

Speaking Elasticsearch: a translation guide from pandas and SQL

IN PANDAS / SQL		IN ELASTICSEARCH	WHAT IT MEANS IN PRACTICE
DataFrame / table	→	Index	A named collection, e.g. ookla-tiles
Row	→	Document	One tile per quarter, stored as JSON with a unique <code>_id</code>
Column	→	Field	quadkey, avg_d_kbps, tile ... each is indexed
dtypes / schema	→	Mapping	Declared types; changing one later means reindexing
df[filter] / WHERE	→	Query	Query DSL (JSON) or ES QL (piped text)
groupby().agg() / GROUP BY	→	Aggregation	Buckets (groups) and metrics (statistics)

Anatomy of a cluster: nodes, shards and replicas



ookla-tiles: 3 primaries (P) × 1 replica (R) — never on the same node.

Source: Elastic documentation, "Size your shards".



Node

One running Elasticsearch process, usually one machine.



Shard

A slice of an index; the unit that moves between nodes.



Replica

A copy on another node: resilience plus extra search capacity.

SIZING GUIDANCE

10–50 GB

per shard, with fewer than 200 million documents each. A national Ookla index needs only one or two primaries.

SECTION 03

03

Modelling statistical data

How do we turn an Ookla tile into a searchable document?



From a Parquet row to an enriched JSON document

AS PUBLISHED BY OOKLA

```
quadkey      "3001010303300022"  
avg_d_kbps   21450  
avg_u_kbps   9870  
avg_lat_ms   38  
tests        142  
devices      57  
tile         "POLYGON((30.0586  
              -1.9552, ...))"
```

ENRICH

INDEXED IN ELASTICSEARCH

```
{  
  "quadkey": "3001010303300022",  
  "service_type": "mobile",  
  "quarter": "2025-Q2",  
  "quarter_start": "2025-04-01",  
  "avg_d_kbps": 21450, "avg_u_kbps": 9870,  
  "avg_lat_ms": 38, "tests": 142, "devices": 57,  
  "admin1": "Kigali City",  
  "admin2": "Nyarugenge",  
  "pop_worldpop": 4830.5,  
  "centroid": {"lat": -1.9525, "lon": 30.0614},  
  "tile": "POLYGON((30.0586 -1.9552, ...))"  
}
```

Green text = added during enrichment

- **Quarter & service** from the file path
- **admin1 / admin2** spatial join with boundaries
- **pop_worldpop** zonal sum from the morning
- **centroid** computed from the polygon

Quadkey and polygon are a real zoom-16 tile in Kigali; speeds, tests and population are illustrative values.

Mapping: declare the schema before you index

```
PUT ookla-tiles
{
  "mappings": {
    "dynamic": "strict",  ◀ reject surprises
    "properties": {
      "quadkey": { "type": "keyword"},
      "service_type": { "type": "keyword"},
      "quarter": { "type": "keyword"},
      "quarter_start": { "type": "date"},
      "avg_d_kbps": { "type": "integer"},
      "avg_u_kbps": { "type": "integer"},
      "avg_lat_ms": { "type": "integer"},
      "tests": { "type": "integer"},
      "devices": { "type": "integer"},
      "admin1": { "type": "keyword"},
      "admin2": { "type": "keyword"},
      "pop_worldpop": { "type": "float"},
      "centroid": { "type": "geo_point"},
      "tile": { "type": "geo_shape" }
    }
  }
}
```

keyword

Exact values: codes, names, categories.
Filters, grouping, sorting.

integer · float

Measurements and counts. Ranges and statistics.

date

Time filters and date histograms
(quarter_start).


geo_point · geo_shape

Coordinates and polygons for spatial queries and maps.



Keep codes as text. A tile over Dakar has quadkey 0333...; cast to a number, it silently loses its leading zero. Map quadkeys and admin codes as keyword.


geo_point or geo_shape? Two views of the same tile



geo_point

the tile centroid

- Compact and very fast
- Heatmaps and geotile_grid buckets
- Distance queries ("within 10 km")
- A border tile belongs to one side only



geo_shape

the full 610 m polygon

- Exact intersection with boundaries
- Border tiles counted by both regions
- Accurate polygon rendering in Maps
- Heavier to index and to query



Recommendation: store both. Centroid for speed, polygon for exact boundary intersections. Mind the axis order: WKT and GeoJSON are longitude first, as in the Ookla files.

Four ways to get data in — and the one we use today



LAB 1

Python client

Bulk helper in a notebook.
Scriptable, versioned, reproducible.



Kibana file upload

Drag-and-drop a small CSV or
NDJSON file to explore quickly.



Logstash

Recurring pipelines: input, parse,
enrich, output.



Ingest pipelines

Processing inside Elasticsearch:
convert, rename, enrich.

```
import pandas as pd
from elasticsearch import Elasticsearch
from elasticsearch import helpers

es = Elasticsearch(URL, api_key=KEY)
df = pd.read_parquet("rw_tiles.parquet")

def actions(df):
    for r in df.to_dict("records"):
        yield {
            "_index": "ookla-tiles",
            # same tile = same id: no duplicates
            "_id": f"{r['quadkey']}-{r['quarter']}-mobile",
            "_source": r}

ok, errors = helpers.bulk(
    es, actions(df), chunk_size=5000)
```

SECTION 04

04

Querying and aggregating

How do we ask the cluster statistical questions?



Anatomy of a search request

```
GET ookla-tiles/_search
{
  "size": 0,      ◀ 1
  "query": {
    "bool": {
      "filter": [ ◀ 2
        { "term": { "service_type": "mobile" } },
        { "term": { "quarter": "2025-Q2" } }
      ]
    }
  },
  "aggs": {
    "by_region": {
      "terms": { "field": "admin1", "size": 50 }, ◀ 3
      "aggs": {
        "mean_dl": {
          "weighted_avg": { ◀ 4
            "value": { "field": "avg_d_kbps" },
            "weight": { "field": "tests" }
          }
        }
      }
    }
  }
}
```

1

size: 0

Return only the aggregated results — not the thousands of documents behind them.

2

filter

Yes/no conditions: no relevance scoring, cacheable, fast. This is your WHERE clause.

3

terms

One bucket per region — the GROUP BY. Set size above your number of regions.

4

weighted_avg

Mean download speed per region, weighted by the number of tests in each tile.

The aggregation toolbox — and its statistical fine print



Metric

compute a value

`avg · sum · min · max`

≈ mean(), sum()

`weighted_avg`

≈ np.average(w=)

`percentiles`

≈ quantile()

`cardinality`

≈ nunique()



Bucket

build groups

`terms`

≈ groupby(col)

`range · histogram`

≈ pd.cut()

`date_histogram`

≈ resample()

`geotile_grid`

≈ group by map tile



Pipeline

work on other aggs

`derivative`

≈ diff()

`cumulative_sum`

≈ cumsum()

`bucket_sort`

≈ sort + head()

`bucket_script`

≈ computed ratio



Approximate by design: cardinality (HyperLogLog++), percentiles (TDigest) and multi-shard terms counts. Fine for exploration — verify precision or compute exactly before publishing.


Building a population-weighted connectivity indicator

FORMULA FOR REGION r

$$C_r = \Sigma (\text{pop}_i \times \text{speed}_i) \div \Sigma \text{pop}_i$$

Tile	Population	Tests	Mbps
A · urban	5,000	400	20
B · peri-urban	1,200	40	8
C · rural	300	5	2


Same three tiles, three answers (Mbps)



Two rules. Document which weight you used — it changes the result. And tiles without tests have no speed, so their population drops out: publish the share of population covered next to the indicator.

Illustrative figures for teaching purposes, not observed data. In Elasticsearch: `weighted_avg` with weight field `pop_worldpop`.


Asking geographic questions



geo_shape

Which tiles intersect this district?


Subnational indicators, border tiles included.



geo_distance

Which tiles lie within 10 km of this point?

Access around hospitals, schools or markets.



geotile_grid

How many tests per map cell?

Heatmaps and multi-scale summaries.

REUSE A STORED BOUNDARY INSTEAD OF PASTING COORDINATES

```
"query": {"bool": {"filter": {"geo_shape": {"tile": {  
  "indexed_shape": {"index": "admin-boundaries", "id": "RW-02", "path": "geometry"},  
  "relation": "intersects"}}}}}}
```

ES|QL: the same question, written as a pipeline

```
FROM ookla-tiles
| WHERE service_type == "mobile"
  AND quarter == "2025-Q2"
| EVAL w = TO_DOUBLE(avg_d_kbps) * tests ◀ use decimals
| STATS sum_w = SUM(w),
      n_tests = SUM(tests) BY admin1
| EVAL mean_mbps = ROUND(sum_w / n_tests / 1000, 1)
| SORT mean_mbps DESC
| KEEP admin1, mean_mbps, n_tests
```



Each step receives the table produced by the previous one — like a pipe in R or a chain in pandas.

Source: Elastic, "Elastic 8.14: Announcing the GA of ES|QL" (June 2024).

Query DSL · JSON

- Every Elasticsearch feature
- Best for applications and dashboards
- Verbose, but precise

ES|QL · piped text

- Reads like a recipe
- Ideal for exploration in Discover
- GA since 8.14 (June 2024)

SECTION 05

05

Exploring with Kibana

How do colleagues who do not code get answers from the cluster?



Kibana: five tools you will use this afternoon



Dev Tools

Console to run Query DSL and ES|QL requests.

Lab 1: mapping, indexing checks, first queries



Discover

Browse and filter documents; inspect fields.

Lab 2: sanity checks, KQL filters



Lens

Drag-and-drop charts: bars, lines, tables.

Lab 2: speed by quarter, fixed vs mobile



Maps

Tile polygons, points and choropleths.

Lab 2: region map, tile zoom-in



Dashboards


Combine panels with shared filters.

Lab 2: the connectivity explorer

KQL — THE FILTER BAR LANGUAGE IN DISCOVER AND DASHBOARDS

```
service_type : "mobile" and tests >= 5 and admin1 : "Kigali City"
```

Lab 2 workflow: building a search-driven connectivity explorer



Deliverable: a dashboard that answers "Where is mobile connectivity weakest, and is it improving?" — plus a screenshot and the saved-objects export (NDJSON) committed to your team repository.

pandas or Elasticsearch? Decide with your own evidence



Stay with pandas / DuckDB when...

- The data fits comfortably in memory
- One analyst, one-off or periodic analysis
- The reproducible script is the product
- No server skills available to maintain a cluster



Move to Elasticsearch when...

- Data outgrows a single machine
- Many colleagues need interactive answers
- Repeated filtered and spatial queries
- New data arrives continuously

YOUR BENCHMARK — FILL IN DURING LAB 2

Task (same data, same question)	pandas time	pandas memory	ES query time	ES disk / heap
1 · Load or index one quarter	...	...	...	...
2 · Weighted mean speed by region	...	...	...	...
3 · Tiles intersecting one district	...	...	...	...

SECTION 06

06

Governing and sustaining

What must be true before an NSO publishes from a search engine?



Publishing on GitHub: the repository your team leaves with

SUGGESTED STRUCTURE

```
stg17-<country>-connectivity/  
├── README.md  
├── LICENSE  
├── mappings/  
│   └── ookla-tiles.json  
├── notebooks/  
│   ├── 01_index_tiles.ipynb  
│   └── 02_benchmark.ipynb  
├── queries/  
│   ├── by_region.json  
│   └── by_region.esql  
├── kibana/  
│   └── dashboard.ndjson  
└── outputs/  
    └── connectivity_by_region.csv
```

THE README MUST STATE



Method

Enrichment, weighting choice, minimum-tests threshold.



Sources

Ookla, WorldPop and boundaries — with vintage and access date.



Licences

CC BY-NC-SA 4.0 for derived data; a separate licence for code.



Limitations

Coverage bias, population not covered, what the proxy cannot support.



How to reproduce

Software versions, commands, expected outputs.

Feeds → Action Plan 2.1.1 · Day 5 country presentations

Six ideas to take back to your office

01

Search is one block of the chain

Keep raw files in the lake; index a curated, enriched subset.

02

The inverted index is the secret

Work done at indexing time buys speed at query time.

03

Model before you load

Strict mappings, codes as keyword, both point and shape.

04

Weights are a methodological choice

Document them, and publish population coverage.

05

Kibana opens data to non-coders

But decide on tooling with your own benchmark.

06

Speed never replaces quality

Licence, confidentiality and honest limitations come first.

Sources

	Workshop context	AfDB, AU STATAFRIC & NISR — STG17 Technical Workshop "Emerging Issues, Emerging Practice": agenda and STG17 Action Plan 2025–2030 links (v. 01/09/2026).
	Ookla data	Ookla — ookla-open-data README, github.com/teamookla/ookla-open-data (tile size, fields, quarterly releases since Q1 2019).
	Ookla licence	Registry of Open Data on AWS — Speedtest by Ookla Global Fixed and Mobile Network Performance Maps (CC BY-NC-SA 4.0).
	Shard sizing	Elastic documentation — "Size your shards", elastic.co/docs (10–50 GB per shard, under 200 million documents).
	ES QL	Elastic blog — "Elastic 8.14: Announcing the GA of Elasticsearch Query Language (ES QL)", June 2024.
	Licensing	Elastic — "FAQ on Software Licensing" and S. Banon, "Elasticsearch is Open Source. Again!" (AGPLv3 option, 2024).
	OpenSearch	Reverera blog — "Elastic's Return to Open Source" (January 2025): OpenSearch fork and transfer to the Linux Foundation.

QUESTIONS · THEN TO THE LAB

From tiles to insight — at scale

Your cluster is running. Your notebooks are waiting. Sit with your team.

14:45

Lab 1 — Ookla at scale with Elasticsearch

15:45

Lab 2 — Search-driven exploration

16:45

Commit to the workshop GitHub organisation

African Development Bank · African Union (STATAFRIC) · National Institute of Statistics of Rwanda